

一、进程原理

1、进程

Linux内核把进程称为**任务(task)**，进程的虚拟地址空间分为用户虚拟地址空间和内核虚拟地址空间，所有进程共享内核虚拟地址空间，每个进程有独立的用户空间虚拟地址空间。

进程有两种特殊形式：**没有用户虚拟地址空间的进程称为内核线程，共享用户虚拟地址空间的进程称为用户线程。**通用在不会引起混淆的情况下把用户线程简称为**线程**。共享同一个用户虚拟地址空间的所有用户线程组成一个线程组。

C标准库的进程专业术语和Linux内核的进程专业术语对应关系如下：

C标准库的进程专业术语	Linux内核的
包含多个线程的进程	线程组
只有一个线程的进程	进程或任务
线程	共享用户虚拟地址空间的进

2、Linux进程四要素：

- a.有一段程序供其执行。
- b.有进程专用的系统堆栈空间。
- c.在内核有task_struct数据结构；
- d.有独立的存储空间，拥有专有的用户空间。

如果只具备前三条而缺少第四条，则称为“**线程**”。如果完全没有用户空间，就称为“**内核线程**”；而如果共享用户空间映射就称为“**用户线程**”。内核为每个进程分配一个task_struct结构时。实际分配两个连续物理页面(8192字节)，数据结构task_struct的大小约占1kb字节左右，进程的系统空间堆栈的大小约为7kb字节（不能扩展，是静态确定的）。

3、进程描述符task_struct数据结构内核源码，其主要核心成员如下：

```
include > linux > C sched.h > task_struct
484
485 struct task_struct {
486 #ifdef CONFIG_THREAD_INFO_IN_TASK
487     /*
488      * For reasons of header soup (see current_thread_info()), this
489      * must be the first element of task_struct.
490      */
491     struct thread_info    thread_info;
492 #endif
493     /* -1 unrunnable, 0 runnable, >0 stopped: */
494     volatile long        state;
495     void                  *stack;
496     atomic_t              usage;
497     /* Per task flags (PF_*), defined further below: */
498     unsigned int          flags;
499     unsigned int          ptrace;
500
```

```
struct task_struct {
```

```
#ifdef CONFIG_THREAD_INFO_IN_TASK
```

```
/*
```

```
 * For reasons of header soup (see current_thread_info()), this
 * must be the first element of task_struct.
```

```
*/
```

```
struct thread_info    thread_info;
```

```
#endif
```

```
/* -1 unrunnable, 0 runnable, >0 stopped: */
```

```
volatile long        state;
```

```
void                  *stack;
```

```
atomic_t              usage;
```

```
/* Per task flags (PF_*), defined further below: */
```

```
unsigned int          flags;
```

```
unsigned int          ptrace;
```

```
#ifdef CONFIG_SMP
```

```
struct llist_node     wake_entry;
```

```
int            on_cpu;

#ifdef CONFIG_THREAD_INFO_IN_TASK
/* Current CPU: */
unsigned int    cpu;
#endif

unsigned int    wakee_flips;
unsigned long    wakee_flip_decay_ts;
struct task_struct *last_wakee;
int            wake_cpu;
#endif

int            on_rq;
int            prio;
int            static_prio;
int            normal_prio;
unsigned int     rt_priority;
const struct sched_class *sched_class;
struct sched_entity se;
struct sched_rt_entity rt;
#ifdef CONFIG_CGROUP_SCHED
struct task_group *sched_task_group;
#endif

struct sched_dl_entity dl;
#ifdef CONFIG_PREEMPT_NOTIFIERS
/* List of struct preempt_notifier: */
struct hlist_head preempt_notifiers;
#endif

#ifdef CONFIG_BLK_DEV_IO_TRACE
```

```

    unsigned int          btrace_seq;
#endif
    unsigned int          policy;
    int                   nr_cpus_allowed;
    cpumask_t             cpus_allowed;
#ifdef CONFIG_PREEMPT_RCU
    int                   rcu_read_lock_nesting;
    union rcu_special     rcu_read_unlock_special;
    struct list_head      rcu_node_entry;
    struct rcu_node       *rcu_blocked_node;
#endif /* #ifdef CONFIG_PREEMPT_RCU */
#ifdef CONFIG_TASKS_RCU
    unsigned long         rcu_tasks_nvcsw;
    bool                  rcu_tasks_holdout;
    struct list_head      rcu_tasks_holdout_list;
    int                   rcu_tasks_idle_cpu;
#endif /* #ifdef CONFIG_TASKS_RCU */
    struct sched_info     sched_info;
    struct list_head      tasks;

```

// 配置SMP（SMP（对称多处理器）系统的应用越来越广泛，规模也越来越大，但由于传统的SMP系统中，所有处理器都共享系统总线，因此当处理器的数目增大时，系统总线的竞争冲突加大，系统总线将成为瓶颈，所以目前SMP系统的CPU数目一般只有数十个，可扩展能力受到极大限制。NUMA技术有效结合了SMP系统易编程性和MPP（大规模并行）系统易扩展性的特点，较好解决了SMP系统的可扩展性问题，已成为当今高性能服务器的主流体系结构之一。

```
#ifdef CONFIG_SMP
```

```
    struct plist_node    pushable_tasks;  
    struct rb_node       pushable_dl_tasks;
```

```
#endif
```

```
    struct mm_struct     *mm;  
    struct mm_struct     *active_mm;
```

```
/* Per-thread vma caching: */
```

```
    struct vmacache      vmacache;
```

```
#ifdef SPLIT_RSS_COUNTING
```

```
    struct task_rss_stat  rss_stat;
```

```
#endif
```

```
    int                  exit_state;
```

```
    int                  exit_code;
```

```
    int                  exit_signal;
```

```
/* The signal sent when the parent dies: */
```

```
    int                  pdeath_signal;
```

```
/* JOBCTL_*, siglock protected: */
```

```
    unsigned long        jobctl;
```

```
/* Used for emulating ABI behavior of previous Linux versions: */
```

```
    unsigned int         personality;
```

```
/* Scheduler bits, serialized by scheduler locks: */
```

```
    unsigned             sched_reset_on_fork:1;
```

```
    unsigned             sched_contributes_to_load:1;
```

```
    unsigned             sched_migrated:1;
```

```
    unsigned             sched_remote_wakeup:1;
```

```
/* Force alignment to the next boundary: */
```

```
    unsigned             :0;
```

```

/* Unserialized, strictly 'current' */
/* Bit to tell LSMs we're in execve(): */
unsigned      in_execve:1;
unsigned      in_iowait:1;
#ifdef TIF_RESTORE_SIGMASK
    unsigned      restore_sigmask:1;
#endif
#ifdef CONFIG_MEMCG
    unsigned      memcg_may_oom:1;
#ifdef CONFIG_SLOB
    unsigned      memcg_kmem_skip_account:1;
#endif
#endif
#ifdef CONFIG_COMPAT_BRK
    unsigned      brk_randomized:1;
#endif
#ifdef CONFIG_CGROUPS
    /* disallow userland-initiated cgroup migration */
    unsigned      no_cgroup_migration:1;
#endif
    unsigned long      atomic_flags; /* Flags requiring atomic access. */
    struct restart_block      restart_block;
    pid_t      pid;
    pid_t      tgid;
#ifdef CONFIG_CC_STACKPROTECTOR
    /* Canary value for the -fstack-protector GCC feature: */

```

```

    unsigned long        stack_canary;
#endif
    /*
     * Pointers to the (original) parent process, youngest child, young
er sibling,
     * older sibling, respectively. (p->father can be replaced with
     * p->real_parent->pid)
     */
    /* Real parent process: */
    struct task_struct __rcu *real_parent;
    /* Recipient of SIGCHLD, wait4() reports: */
    struct task_struct __rcu *parent;
    /*
     * Children/sibling form the list of natural children:
     */
    struct list_head      children;
    struct list_head      sibling;
    struct task_struct    *group_leader;
    /*
     * 'ptraced' is the list of tasks this task is using ptrace() on.
     *
     * This includes both natural children and PTRACE_ATTACH target
s.
     * 'ptrace_entry' is this task's link on the p->parent->ptraced list.
     */
    struct list_head      ptraced;
    struct list_head      ptrace_entry;

```

```

/* PID/PID hash table linkage. */
struct pid_link      pids[PIDTYPE_MAX];
struct list_head      thread_group;
struct list_head      thread_node;
struct completion      *vfork_done;
/* CLONE_CHILD_SETTID: */
int __user      *set_child_tid;
/* CLONE_CHILD_CLEARTID: */
int __user      *clear_child_tid;
u64      utime;
u64      stime;
#ifdef CONFIG_ARCH_HAS_SCALED_CPUTIME
u64      utimescaled;
u64      stimescaled;
#endif
u64      gtime;
struct prev_cputime      prev_cputime;
#ifdef CONFIG_VIRT_CPU_ACCOUNTING_GEN
seqcount_t      vtime_seqcount;
unsigned long long      vtime_snap;
enum {
    /* Task is sleeping or running in a CPU with VTIME inactive: */
    VTIME_INACTIVE = 0,
    /* Task runs in userspace in a CPU with VTIME active: */
    VTIME_USER,
    /* Task runs in kernelspace in a CPU with VTIME active: */
    VTIME_SYS,

```



```

    } vtime_snap_whence;
#endif
#ifdef CONFIG_NO_HZ_FULL
    atomic_t      tick_dep_mask;
#endif

    /* Context switch counts: */
    unsigned long    nvcs;
    unsigned long    nivcs;
    /* Monotonic time in nsecs: */
    u64      start_time;
    /* Boot based time in nsecs: */
    u64      real_start_time;
    /* MM fault and swap info: this can arguably be seen as either m
m-specific or thread-specific: */
    unsigned long    minflt;
    unsigned long    majflt;
#ifdef CONFIG_POSIX_TIMERS
    struct task_cputime    cputime_expires;
    struct list_head    cpu_timers[3];
#endif

    /* Process credentials: */
    /* Tracer's credentials at attach: */
    const struct cred __rcu    *ptracer_cred;
    /* Objective and real subjective task credentials (COW): */
    const struct cred __rcu    *real_cred;
    /* Effective (overridable) subjective task credentials (COW): */
    const struct cred __rcu    *cred;

```

```

/*
 * executable name, excluding path.
 *
 * - normally initialized setup_new_exec()
 * - access it with [gs]et_task_comm()
 * - lock it with task_lock()
 */
char          comm[TASK_COMM_LEN];
struct nameidata *nameidata;
#ifdef CONFIG_SYSVIPC
    struct sysv_sem    sysvsem;
    struct sysv_shm    sysvshm;
#endif
#ifdef CONFIG_DETECT_HUNG_TASK
    unsigned long      last_switch_count;
#endif
/* Filesystem information: */
struct fs_struct      *fs;
/* Open file information: */
struct files_struct   *files;
/* Namespaces: */
struct nsproxy        *nsproxy;
/* Signal handlers: */
struct signal_struct   *signal;
struct sighand_struct  *sighand;
sigset_t              blocked;
sigset_t              real_blocked;

```

```

/* Restored if set_restore_sigmask() was used: */
sigset_t      saved_sigmask;
struct sigpending pending;
unsigned long  sas_ss_sp;
size_t        sas_ss_size;
unsigned int   sas_ss_flags;
struct callback_head *task_works;
struct audit_context *audit_context;
#ifdef CONFIG_AUDITSYSCALL
    kuid_t      loginuid;
    unsigned int sessionid;
#endif

struct seccomp      seccomp;
/* Thread group tracking: */
u32      parent_exec_id;
u32      self_exec_id;

/* Protection against (de-)allocation: mm, files, fs, tty, keyrings, m
ems_allowed, mempolicy: */
spinlock_t      alloc_lock;
/* Protection of the PI data structures: */
raw_spinlock_t      pi_lock;
struct wake_q_node      wake_q;
#ifdef CONFIG_RT_MUTEXES
/* PI waiters blocked on a rt_mutex held by this task: */
struct rb_root      pi_waiters;
struct rb_node      *pi_waiters_leftmost;
/* Updated under owner's pi_lock and rq lock */

```

```

struct task_struct    *pi_top_task;

/* Deadlock detection and priority inheritance handling: */
struct rt_mutex_waiter    *pi_blocked_on;

#endif

#ifdef CONFIG_DEBUG_MUTEXES
/* Mutex deadlock detection: */
struct mutex_waiter    *blocked_on;
#endif

#ifdef CONFIG_TRACE_IRQFLAGS
unsigned int          irq_events;
unsigned long         hardirq_enable_ip;
unsigned long         hardirq_disable_ip;
unsigned int          hardirq_enable_event;
unsigned int          hardirq_disable_event;
int                   hardirqs_enabled;
int                   hardirq_context;
unsigned long         softirq_disable_ip;
unsigned long         softirq_enable_ip;
unsigned int          softirq_disable_event;
unsigned int          softirq_enable_event;
int                   softirqs_enabled;
int                   softirq_context;
#endif

#ifdef CONFIG_LOCKDEP
# define MAX_LOCK_DEPTH    48UL
u64                   curr_chain_key;
int                   lockdep_depth;

```

```

    unsigned int          lockdep_recursion;
    struct held_lock       held_locks[MAX_LOCK_DEPTH];
    gfp_t                  lockdep_reclaim_gfp;
#endif

#ifdef CONFIG_UBSAN
    unsigned int          in_ubsan;
#endif

    /* Journalling filesystem info: */
    void                  *journal_info;
    /* Stacked block device info: */
    struct bio_list        *bio_list;
#ifdef CONFIG_BLOCK
    /* Stack plugging: */
    struct blk_plug        *plug;
#endif

    /* VM state: */
    struct reclaim_state   *reclaim_state;
    struct backing_dev_info *backing_dev_info;
    struct io_context      *io_context;
    /* Ptrace state: */
    unsigned long          ptrace_message;
    siginfo_t              *last_siginfo;
    struct task_io_accounting ioac;
#ifdef CONFIG_TASK_XACCT
    /* Accumulated RSS usage: */
    u64                    acct_rss_mem1;
    /* Accumulated virtual memory usage: */

```

```

    u64          acct_vm_mem1;
    /* stime + utime since last update: */
    u64          acct_timexpd;
#endif

#ifdef CONFIG_CPUSETS
    /* Protected by ->alloc_lock: */
    nodemask_t    mems_allowed;
    /* Sequence number to catch updates: */
    seqcount_t    mems_allowed_seq;
    int           cpuset_mem_spread_rotor;
    int           cpuset_slab_spread_rotor;
#endif

#ifdef CONFIG_CGROUPS
    /* Control Group info protected by css_set_lock: */
    struct css_set __rcu    *cgroups;
    /* cg_list protected by css_set_lock and tsk->alloc_lock: */
    struct list_head    cg_list;
#endif

#ifdef CONFIG_INTEL_RDT_A
    int             closid;
#endif

#ifdef CONFIG_FUTEX
    struct robust_list_head __user *robust_list;
#endif

#ifdef CONFIG_COMPAT
    struct compat_robust_list_head __user *compat_robust_list;
#endif

    struct list_head    pi_state_list;

```

```
    struct futex_pi_state    *pi_state_cache;
#endif
#ifdef CONFIG_PERF_EVENTS
    struct perf_event_context *perf_event_ctxp[perf_nr_task_contexts];

    struct mutex    perf_event_mutex;
    struct list_head perf_event_list;
#endif
#ifdef CONFIG_DEBUG_PREEMPT
    unsigned long    preempt_disable_ip;
#endif
#ifdef CONFIG_NUMA
    /* Protected by alloc_lock: */
    struct mempolicy *mempolicy;
    short    il_next;
    short    pref_node_fork;
#endif
#ifdef CONFIG_NUMA_BALANCING
    int    numa_scan_seq;
    unsigned int    numa_scan_period;
    unsigned int    numa_scan_period_max;
    int    numa_preferred_nid;
    unsigned long    numa_migrate_retry;
    /* Migration stamp: */
    u64    node_stamp;
    u64    last_task_numa_placement;
    u64    last_sum_exec_runtime;
```

```

struct callback_head      numa_work;
struct list_head          numa_entry;
struct numa_group         *numa_group;
/*
 * numa_faults is an array split into four regions:
 * faults_memory, faults_cpu, faults_memory_buffer, faults_cpu_b
uffer
 * in this precise order.
 *
 * faults_memory: Exponential decaying average of faults on a per
-node
 * basis. Scheduling placement decisions are made based on thes
e
 * counts. The values remain static for the duration of a PTE scan.
 * faults_cpu: Track the nodes the process was running on when a
NUMA
 * hinting fault was incurred.
 * faults_memory_buffer and faults_cpu_buffer: Record faults per
node
 * during the current scan window. When the scan completes, the
counts
 * in faults_memory and faults_cpu decay and these values are co
pied.
 */
unsigned long              *numa_faults;
unsigned long              total_numa_faults;
/*

```



```

    * numa_faults_locality tracks if faults recorded during the last
    * scan window were remote/local or failed to migrate. The task s
can
    * period is adapted based on the locality of the faults with differ
ent
    * weights depending on whether they were shared or private fau
lts
    */
    unsigned long        numa_faults_locality[3];
    unsigned long        numa_pages_migrated;
#endif /* CONFIG_NUMA_BALANCING */
    struct tlbflush_unmap_batch tlb_ubic;
    struct rcu_head        rcu;
    /* Cache last used pipe for splice(): */
    struct pipe_inode_info    *splice_pipe;
    struct page_frag        task_frag;
#ifdef CONFIG_TASK_DELAY_ACCT
    struct task_delay_info    *delays;
#endif
#ifdef CONFIG_FAULT_INJECTION
    int        make_it_fail;
#endif
    /*
    * When (nr_dirtied >= nr_dirtied_pause), it's time to call
    * balance_dirty_pages() for a dirty throttling pause:
    */
    int        nr_dirtied;

```

```

int          nr_dirtied_pause;
/* Start of a write-and-pause period: */
unsigned long    dirty_paused_when;
#ifdef CONFIG_LATENCYTOP
int          latency_record_count;
struct latency_record    latency_record[LT_SAVECOUNT];
#endif
/*
 * Time slack values; these are used to round up poll() and
 * select() etc timeout values. These are in nanoseconds.
 */
u64          timer_slack_ns;
u64          default_timer_slack_ns;
#ifdef CONFIG_KASAN
unsigned int    kasan_depth;
#endif
#ifdef CONFIG_FUNCTION_GRAPH_TRACER
/* Index of current stored address in ret_stack: */
int          curr_ret_stack;
/* Stack of return addresses for return function tracing: */
struct ftrace_ret_stack    *ret_stack;
/* Timestamp for last schedule: */
unsigned long long    ftrace_timestamp;
/*
 * Number of functions that haven't been traced
 * because of depth overrun:
 */

```

```

    atomic_t          trace_overrun;
    /* Pause tracing: */
    atomic_t          tracing_graph_pause;
#endif

#ifdef CONFIG_TRACING
    /* State flags for use by tracers: */
    unsigned long      trace;
    /* Bitmask and counter of trace recursion: */
    unsigned long      trace_recursion;
#endif /* CONFIG_TRACING */

#ifdef CONFIG_KCOV
    /* Coverage collection mode enabled for this task (0 if disabled):
    */
    enum kcov_mode      kcov_mode;
    /* Size of the kcov_area: */
    unsigned int        kcov_size;
    /* Buffer for coverage collection: */
    void                *kcov_area;
    /* KCOV descriptor wired with this task or NULL: */
    struct kcov          *kcov;
#endif

#ifdef CONFIG_MEMCG
    struct mem_cgroup    *memcg_in_oom;
    gfp_t                memcg_oom_gfp_mask;
    int                  memcg_oom_order;
    /* Number of pages to reclaim on returning to userland: */
    unsigned int         memcg_nr_pages_over_high;

```

```
#endif

#ifdef CONFIG_UPROBES
    struct uprobe_task    *utask;
#endif

#if defined(CONFIG_BCACHE) || defined(CONFIG_BCACHE_MODULE)
)

    unsigned int          sequential_io;
    unsigned int          sequential_io_avg;
#endif

#ifdef CONFIG_DEBUG_ATOMIC_SLEEP
    unsigned long          task_state_change;
#endif

    int                   pagefault_disabled;
#endif CONFIG_MMU

    struct task_struct     *oom_reaper_list;
#endif

#ifdef CONFIG_VMAP_STACK
    struct vm_struct       *stack_vm_area;
#endif

#ifdef CONFIG_THREAD_INFO_IN_TASK
    /* A live task holds one reference: */
    atomic_t               stack_refcount;
#endif

#ifdef CONFIG_LIVEPATCH
    int patch_state;
#endif

#ifdef CONFIG_SECURITY
```

```

/* Used by LSM modules for access restriction: */
void                *security;
#endif

/* CPU-specific state of this task: */
struct thread_struct    thread;

/*
 * WARNING: on x86, 'thread_struct' contains a variable-sized
 * structure. It *MUST* be at the end of 'task_struct'.
 *
 * Do not put anything below here!
 */
};

```

4、创建新进程

在Linux内核中，新进程是从一个已经存在的进程复制出来的，内核使用静态数据结构造出0号内核线程，0号内核线程分叉生成1号内核线程和2号内核线程（kthreadd线程）。1号内核线程完成初始化以后装载用户程序，变成1号进程，其他进程都是1号进程或者它的子孙进程分叉生成的；其他内核线程是kthreadd线程分叉生成的。

3个系统调用可以用来创建新的进程：

a.fork(分叉)：子进程是父进程的一个副本，采用定时复制技术。

b.vfor：用于创建子进程，之后子进程立即调用execve以装载新程序的情况，为了避免复制物理页，父进程会睡眠等待子进程装载新程序。现在fork采用了定时复制技术，vfork失去了速度优势，已经被废弃。

c.clone（克隆）：可以精确地控制子进程和父进程共享哪些资源。这个系统调用的主要用处是可供pthread库用来创建线程。clone是功能最

齐全的函数，参数多使用复杂，fork是clone的简化函数。

Linux内核定义系统调用的独特方式，目前以系统调用fork为例：

系统调用的函数名称以"sys_"开头，创建新进程的3个系统调用在文件"kernel/fork.c"中，它们把工作委托给函数__do_fork。

```
kernel > C fork.c > ...
2001 |
2002 | long __do_fork(unsigned long clone_flags,
2003 |               unsigned long stack_start,
2004 |               unsigned long stack_size,
2005 |               int __user *parent_tidptr,
2006 |               int __user *child_tidptr,
2007 |               unsigned long tls)
2008 | {
2009 |     struct task_struct *p;
2010 |     int trace = 0;
2011 |     long nr;
2012 |
```

函数__do_fork()内核源码如下：

```
long __do_fork(unsigned long clone_flags,
               unsigned long stack_start,
               unsigned long stack_size,
               int __user *parent_tidptr,
               int __user *child_tidptr,
               unsigned long tls)
```

```
{
```

```
    struct task_struct *p;
```

```
    int trace = 0;
```

```
    long nr;
```

```
    /*
```

```
     * Determine whether and which event to report to ptracer. Whe
```

```
n
```

```
     * called from kernel_thread or CLONE_UNTRACED is explicitly
```

```
     * requested, no event is reported; otherwise, report if the event
```

* for the type of forking is enabled.

*/

```
if (!(clone_flags & CLONE_UNTRACED)) {  
    if (clone_flags & CLONE_VFORK)  
        trace = PTRACE_EVENT_VFORK;  
    else if ((clone_flags & CSIGNAL) != SIGCHLD)  
        trace = PTRACE_EVENT_CLONE;  
    else  
        trace = PTRACE_EVENT_FORK;  
    if (likely(!ptrace_event_enabled(current, trace)))  
        trace = 0;  
}
```

```
p = copy_process(clone_flags, stack_start, stack_size,  
                child_tidptr, NULL, trace, tls, NUMA_NO_NODE);
```

```
add_latent_entropy();
```

/*

* Do this prior waking up the new thread - the thread pointer

* might get invalid after that point, if the thread exits quickly.

*/

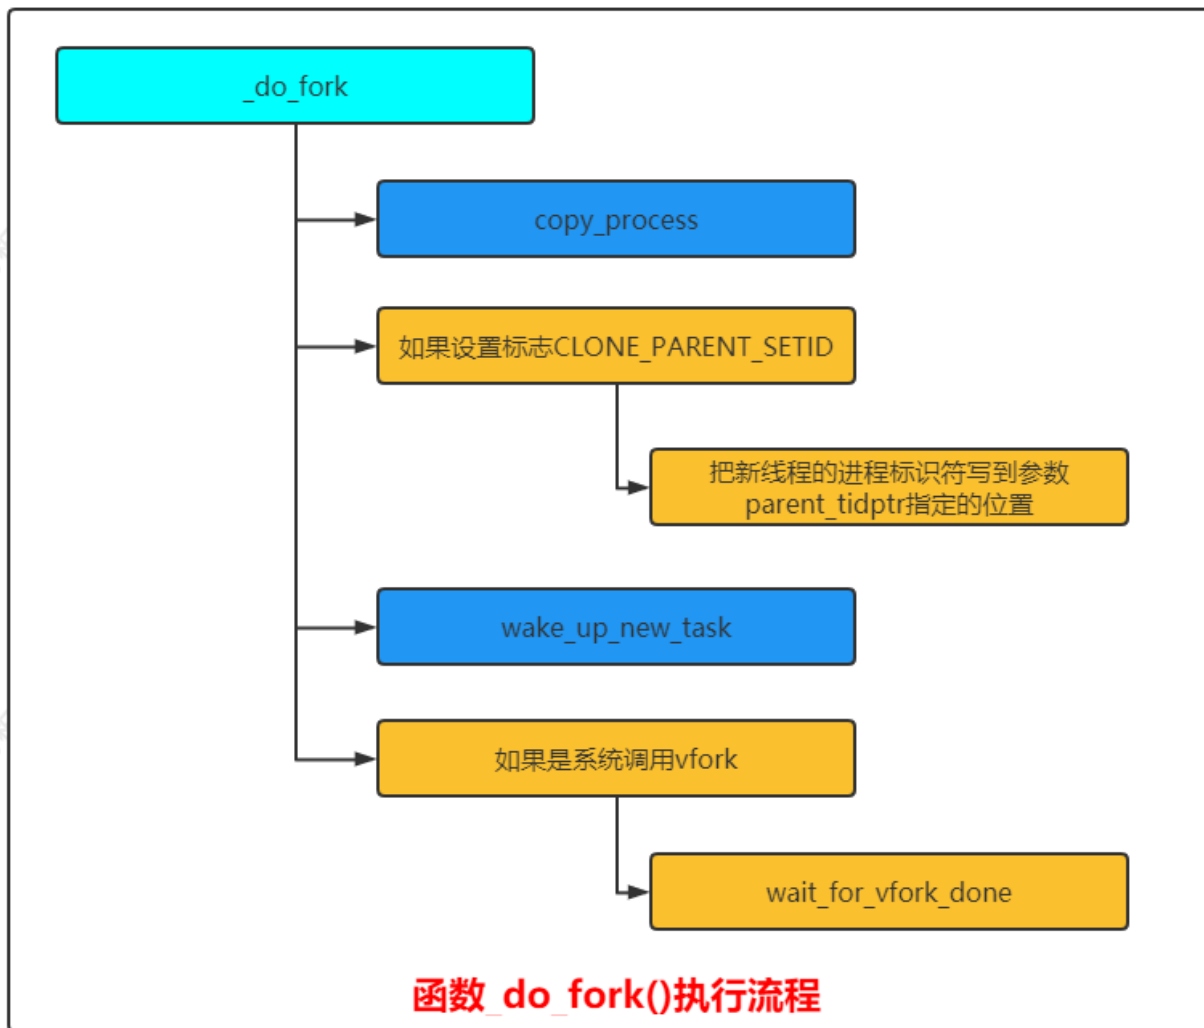
```
if (!IS_ERR(p)) {  
    struct completion vfork;  
    struct pid *pid;  
    trace_sched_process_fork(current, p);  
    pid = get_task_pid(p, PIDTYPE_PID);  
    nr = pid_vnr(pid);  
    if (clone_flags & CLONE_PARENT_SETTID)  
        put_user(nr, parent_tidptr);
```

```

if (clone_flags & CLONE_VFORK) {
    p->vfork_done = &vfork;
    init_completion(&vfork);
    get_task_struct(p);
}
wake_up_new_task(p);
/* forking complete and child started to run, tell ptracer */
if (unlikely(trace))
    ptrace_event_pid(trace, pid);
if (clone_flags & CLONE_VFORK) {
    if (!wait_for_vfork_done(p, &vfork))
        ptrace_event_pid(PTRACE_EVENT_VFORK_DONE, pid);
}
put_pid(pid);
} else {
    nr = PTR_ERR(p);
}
return nr;
}

```

Linux内核函数_do_fork()执行流程如下图所示：



```
kernel > C fork.c > copy_process(unsigned long, unsigned long, unsigned long, int __user *, pid *, int, unsigned long, int)
1516 static __latent_entropy struct task_struct *copy_process({
1517     unsigned long clone_flags,
1518     unsigned long stack_start,
1519     unsigned long stack_size,
1520     int __user *child_tidptr,
1521     struct pid *pid,
1522     int trace,
1523     unsigned long tls,
1524     int node)
1525 {
1526     int retval;
1527     struct task_struct *p;
1528
1529     if ((clone_flags & (CLONE_NEWNS|CLONE_FS)) == (CLONE_NEWNS|CLONE_FS))
1530         return ERR_PTR(-EINVAL);
1531
1532     if ((clone_flags & (CLONE_NEWUSER|CLONE_FS)) == (CLONE_NEWUSER|CLONE_FS))
1533         return ERR_PTR(-EINVAL);
1534 }
```

函数 `copy_process()` 内核源码如下：

```
static __latent_entropy struct task_struct *copy_process(
    unsigned long clone_flags,
    unsigned long stack_start,
```

```

        unsigned long stack_size,
        int __user *child_tidptr,
        struct pid *pid,
        int trace,
        unsigned long tls,
        int node)
{
    int retval;
    struct task_struct *p;
    if ((clone_flags & (CLONE_NEWNS|CLONE_FS)) == (CLONE_NEW
NS|CLONE_FS))
        return ERR_PTR(-EINVAL);
    if ((clone_flags & (CLONE_NEWUSER|CLONE_FS)) == (CLONE_NE
WUSER|CLONE_FS))
        return ERR_PTR(-EINVAL);
    /*
     * Thread groups must share signals as well, and detached thread
s
     * can only be started up within the thread group.
     */
    if ((clone_flags & CLONE_THREAD) && !
(clone_flags & CLONE_SIGHAND))
        return ERR_PTR(-EINVAL);
    /*
     * Shared signal handlers imply shared VM. By way of the above,
     * thread groups also imply shared VM. Blocking this case allows
     * for various simplifications in other code.

```

```

    */
    if ((clone_flags & CLONE_SIGHAND) && !
(clone_flags & CLONE_VM))
        return ERR_PTR(-EINVAL);
    /*
     * Siblings of global init remain as zombies on exit since they are
     * not reaped by their parent (swapper). To solve this and to avoid
     * multi-rooted process trees, prevent global and container-inits
     * from creating siblings.
     */
    if ((clone_flags & CLONE_PARENT) &&
        current->signal->flags & SIGNAL_UNKILLABLE)
        return ERR_PTR(-EINVAL);
    /*
     * If the new process will be in a different pid or user namespace
     * do not allow it to share a thread group with the forking task.
     */
    if (clone_flags & CLONE_THREAD) {
        if ((clone_flags & (CLONE_NEWUSER | CLONE_NEWPID)) ||
            (task_active_pid_ns(current) !=
             current->nsproxy->pid_ns_for_children))
            return ERR_PTR(-EINVAL);
    }
    retval = security_task_create(clone_flags);
    if (retval)
        goto fork_out;

```

```

retval = -ENOMEM;
p = dup_task_struct(current, node);
if (!p)
    goto fork_out;
/*
 * This _must_ happen before we call free_task(), i.e. before we ju
mp
 * to any of the bad_fork_* labels. This is to avoid freeing
 * p-
>set_child_tid which is (ab)used as a kthread's data pointer for
 * kernel threads (PF_KTHREAD).
 */
p-
>set_child_tid = (clone_flags & CLONE_CHILD_SETTID) ? child_tidptr
: NULL;
/*
 * Clear TID on mm_release()?
 */
p-
>clear_child_tid = (clone_flags & CLONE_CHILD_CLEARTID) ? child_t
idptr : NULL;
ftrace_graph_init_task(p);
rt_mutex_init_task(p);
#ifdef CONFIG_PROVE_LOCKING
    DEBUG_LOCKS_WARN_ON(!p->hardirqs_enabled);
    DEBUG_LOCKS_WARN_ON(!p->softirqs_enabled);
#endif

```

```

retval = -EAGAIN;
if (atomic_read(&p->real_cred->user->processes) >=
    task_rlimit(p, RLIMIT_NPROC)) {
    if (p->real_cred->user != INIT_USER &&
        !capable(CAP_SYS_RESOURCE) && !capable(CAP_SYS_ADMIN))
        goto bad_fork_free;
}
current->flags &= ~PF_NPROC_EXCEEDED;
retval = copy_creds(p, clone_flags);
if (retval < 0)
    goto bad_fork_free;
/*
 * If multiple threads are within copy_process(), then this check
 * triggers too late. This doesn't hurt, the check is only there
 * to stop root fork bombs.
 */
retval = -EAGAIN;
if (nr_threads >= max_threads)
    goto bad_fork_cleanup_count;
delayacct_tsk_init(p); /* Must remain after dup_task_struct() */
p->flags &= ~(PF_SUPERPRIV | PF_WQ_WORKER | PF_IDLE);
p->flags |= PF_FORKNOEXEC;
INIT_LIST_HEAD(&p->children);
INIT_LIST_HEAD(&p->sibling);
rcu_copy_process(p);
p->vfork_done = NULL;

```

```
spin_lock_init(&p->alloc_lock);
init_sigpending(&p->pending);
p->utime = p->stime = p->gtime = 0;
#ifdef CONFIG_ARCH_HAS_SCALED_CPUTIME
    p->utimescaled = p->stimescaled = 0;
#endif

prev_cputime_init(&p->prev_cputime);
#ifdef CONFIG_VIRT_CPU_ACCOUNTING_GEN
    seqcount_init(&p->vtime_seqcount);
    p->vtime_snap = 0;
    p->vtime_snap_whence = VTIME_INACTIVE;
#endif

#if defined(SPLIT_RSS_COUNTING)
    memset(&p->rss_stat, 0, sizeof(p->rss_stat));
#endif

p->default_timer_slack_ns = current->timer_slack_ns;
task_io_accounting_init(&p->ioac);
acct_clear_integrals(p);
posix_cpu_timers_init(p);
p->start_time = ktime_get_ns();
p->real_start_time = ktime_get_boot_ns();
p->io_context = NULL;
p->audit_context = NULL;
cgroup_fork(p);
#ifdef CONFIG_NUMA
p->mempolicy = mpol_dup(p->mempolicy);
if (IS_ERR(p->mempolicy)) {
```

```
    retval = PTR_ERR(p->mempolicy);
    p->mempolicy = NULL;
    goto bad_fork_cleanup_threadgroup_lock;
}

#endif

#ifdef CONFIG_CPUSETS
    p->cpuset_mem_spread_rotor = NUMA_NO_NODE;
    p->cpuset_slab_spread_rotor = NUMA_NO_NODE;
    seqcount_init(&p->mems_allowed_seq);
#endif

#ifdef CONFIG_TRACE_IRQFLAGS
    p->irq_events = 0;
    p->hardirqs_enabled = 0;
    p->hardirq_enable_ip = 0;
    p->hardirq_enable_event = 0;
    p->hardirq_disable_ip = _THIS_IP_;
    p->hardirq_disable_event = 0;
    p->softirqs_enabled = 1;
    p->softirq_enable_ip = _THIS_IP_;
    p->softirq_enable_event = 0;
    p->softirq_disable_ip = 0;
    p->softirq_disable_event = 0;
    p->hardirq_context = 0;
    p->softirq_context = 0;
#endif

    p->pagefault_disabled = 0;

#ifdef CONFIG_LOCKDEP
```

```
p->lockdep_depth = 0; /* no locks held yet */
p->curr_chain_key = 0;
p->lockdep_recursion = 0;
#endif

#ifdef CONFIG_DEBUG_MUTEXES
    p->blocked_on = NULL; /* not blocked yet */
#endif

#ifdef CONFIG_BCACHE
    p->sequential_io    = 0;
    p->sequential_io_avg = 0;
#endif

/* Perform scheduler related setup. Assign this task to a CPU. */
retval = sched_fork(clone_flags, p);
if (retval)
    goto bad_fork_cleanup_policy;
retval = perf_event_init_task(p);
if (retval)
    goto bad_fork_cleanup_policy;
retval = audit_alloc(p);
if (retval)
    goto bad_fork_cleanup_perf;
/* copy all the process information */
shm_init_task(p);
retval = security_task_alloc(p, clone_flags);
if (retval)
    goto bad_fork_cleanup_audit;
retval = copy_semundo(clone_flags, p);
```



```
if (retval)
    goto bad_fork_cleanup_security;
retval = copy_files(clone_flags, p);
if (retval)
    goto bad_fork_cleanup_semundo;
retval = copy_fs(clone_flags, p);
if (retval)
    goto bad_fork_cleanup_files;
retval = copy_sighand(clone_flags, p);
if (retval)
    goto bad_fork_cleanup_fs;
retval = copy_signal(clone_flags, p);
if (retval)
    goto bad_fork_cleanup_sighand;
retval = copy_mm(clone_flags, p);
if (retval)
    goto bad_fork_cleanup_signal;
retval = copy_namespaces(clone_flags, p);
if (retval)
    goto bad_fork_cleanup_mm;
retval = copy_io(clone_flags, p);
if (retval)
    goto bad_fork_cleanup_namespaces;
retval = copy_thread_tls(clone_flags, stack_start, stack_size, p, tls);
if (retval)
    goto bad_fork_cleanup_io;
if (pid != &init_struct_pid) {
```

```

    pid = alloc_pid(p->nsproxy->pid_ns_for_children);
    if (IS_ERR(pid)) {
        retval = PTR_ERR(pid);
        goto bad_fork_cleanup_thread;
    }
}

#ifdef CONFIG_BLOCK
    p->plug = NULL;
#endif

#ifdef CONFIG_FUTEX
    p->robust_list = NULL;
#endif

#ifdef CONFIG_COMPAT
    p->compat_robust_list = NULL;
#endif

INIT_LIST_HEAD(&p->pi_state_list);
p->pi_state_cache = NULL;

/*
 * sigaltstack should be cleared when sharing the same VM
 */
if ((clone_flags & (CLONE_VM|CLONE_VFORK)) == CLONE_VM)
    sas_ss_reset(p);

/*
 * Syscall tracing and stepping should be turned off in the
 * child regardless of CLONE_PTRACE.
 */
user_disable_single_step(p);

```

```

    clear_tsk_thread_flag(p, TIF_SYSCALL_TRACE);
#ifdef TIF_SYSCALL_EMU
    clear_tsk_thread_flag(p, TIF_SYSCALL_EMU);
#endif

    clear_all_latency_tracing(p);
    /* ok, now we should be set up.. */
    p->pid = pid_nr(pid);
    if (clone_flags & CLONE_THREAD) {
        p->exit_signal = -1;
        p->group_leader = current->group_leader;
        p->tgid = current->tgid;
    } else {
        if (clone_flags & CLONE_PARENT)
            p->exit_signal = current->group_leader->exit_signal;
        else
            p->exit_signal = (clone_flags & CSIGNAL);
        p->group_leader = p;
        p->tgid = p->pid;
    }
    p->nr_dirtied = 0;
    p->nr_dirtied_pause = 128 >> (PAGE_SHIFT - 10);
    p->dirty_paused_when = 0;
    p->pdeath_signal = 0;
    INIT_LIST_HEAD(&p->thread_group);
    p->task_works = NULL;
    cgroup_threadgroup_change_begin(current);
    /*

```

- * Ensure that the cgroup subsystem policies allow the new process to be
- * forked. It should be noted the the new process's css_set can be changed
- * between here and cgroup_post_fork() if an organisation operation is in progress.

```

*/
retval = cgroup_can_fork(p);
if (retval)
    goto bad_fork_free_pid;
/*
* Make it visible to the rest of the system, but don't wake it up yet.
* Need tasklist lock for parent etc handling!
*/
write_lock_irq(&tasklist_lock);
/* CLONE_PARENT re-uses the old parent */
if (clone_flags & (CLONE_PARENT|CLONE_THREAD)) {
    p->real_parent = current->real_parent;
    p->parent_exec_id = current->parent_exec_id;
} else {
    p->real_parent = current;
    p->parent_exec_id = current->self_exec_id;
}
klp_copy_process(p);
spin_lock(&current->sigband->siglock);

```

```

/*
 * Copy seccomp details explicitly here, in case they were change
d
 * before holding sighand lock.
 */
copy_seccomp(p);
/*
 * Process group and session signals need to be delivered to just
the
 * parent before the fork or both the parent and the child after th
e
 * fork. Restart if a signal comes in before we add the new proces
s to
 * it's process group.
 * A fatal signal pending means that current will exit, so the new
 * thread can't slip out of an OOM kill (or normal SIGKILL).
 */
recalc_sigpending();
if (signal_pending(current)) {
    retval = -ERESTARTNOINTR;
    goto bad_fork_cancel_cgroup;
}
if (unlikely(!(ns_of_pid(pid)-
> nr_hashed & PIDNS_HASH_ADDING))) {
    retval = -ENOMEM;
    goto bad_fork_cancel_cgroup;
}

```

```

if (likely(p->pid)) {
    ptrace_init_task(p, (clone_flags & CLONE_PTRACE) || trace);
    init_task_pid(p, PIDTYPE_PID, pid);
    if (thread_group_leader(p)) {
        init_task_pid(p, PIDTYPE_PGID, task_pgrp(current));
        init_task_pid(p, PIDTYPE_SID, task_session(current));
        if (is_child_reaper(pid)) {
            ns_of_pid(pid)->child_reaper = p;
            p->signal->flags |= SIGNAL_UNKILLABLE;
        }
        p->signal->leader_pid = pid;
        p->signal->tty = tty_kref_get(current->signal->tty);
        /*
         * Inherit has_child_subreaper flag under the same
         * tasklist_lock with adding child to the process tree
         * for propagate_has_child_subreaper optimization.
         */
        p->signal->has_child_subreaper = p->real_parent->signal-
>has_child_subreaper ||
            p->real_parent->signal->is_child_subreaper;
        list_add_tail(&p->sibling, &p->real_parent->children);
        list_add_tail_rcu(&p->tasks, &init_task.tasks);
        attach_pid(p, PIDTYPE_PGID);
        attach_pid(p, PIDTYPE_SID);
        __this_cpu_inc(process_counts);
    } else {
        current->signal->nr_threads++;
    }
}

```

```

    atomic_inc(&current->signal->live);
    atomic_inc(&current->signal->sigcnt);
    list_add_tail_rcu(&p->thread_group,
                     &p->group_leader->thread_group);
    list_add_tail_rcu(&p->thread_node,
                     &p->signal->thread_head);
}
attach_pid(p, PIDTYPE_PID);
nr_threads++;
}
total_forks++;
spin_unlock(&current->sigband->siglock);
syscall_tracepoint_update(p);
write_unlock_irq(&tasklist_lock);
proc_fork_connector(p);
cgroup_post_fork(p);
cgroup_threadgroup_change_end(current);
perf_event_fork(p);
trace_task_newtask(p, clone_flags);
uprobe_copy_process(p, clone_flags);
return p;
bad_fork_cancel_cgroup:
    spin_unlock(&current->sigband->siglock);
    write_unlock_irq(&tasklist_lock);
    cgroup_cancel_fork(p);
bad_fork_free_pid:
    cgroup_threadgroup_change_end(current);

```

```
    if (pid != &init_struct_pid)
        free_pid(pid);
bad_fork_cleanup_thread:
    exit_thread(p);
bad_fork_cleanup_io:
    if (p->io_context)
        exit_io_context(p);
bad_fork_cleanup_namespaces:
    exit_task_namespaces(p);
bad_fork_cleanup_mm:
    if (p->mm)
        mmput(p->mm);
bad_fork_cleanup_signal:
    if (!(clone_flags & CLONE_THREAD))
        free_signal_struct(p->signal);
bad_fork_cleanup_sighand:
    __cleanup_sighand(p->sighand);
bad_fork_cleanup_fs:
    exit_fs(p); /* blocking */
bad_fork_cleanup_files:
    exit_files(p); /* blocking */
bad_fork_cleanup_semundo:
    exit_sem(p);
bad_fork_cleanup_security:
    security_task_free(p);
bad_fork_cleanup_audit:
    audit_free(p);
```



```
bad_fork_cleanup_perf:
```

```
    perf_event_free_task(p);
```

```
bad_fork_cleanup_policy:
```

// 配置NUMA(NUMA(Non Uniform Memory Access)即非一致内存访问架构, 市面上主要有X86_64(JASPER)和MIPS64(XLP)体系。)

```
#ifdef CONFIG_NUMA
```

```
    mpol_put(p->mempolicy);
```

```
bad_fork_cleanup_threadgroup_lock:
```

```
#endif
```

```
    delayacct_tsk_free(p);
```

```
bad_fork_cleanup_count:
```

```
    atomic_dec(&p->cred->user->processes);
```

```
    exit_creds(p);
```

```
bad_fork_free:
```

```
    p->state = TASK_DEAD;
```

```
    put_task_stack(p);
```

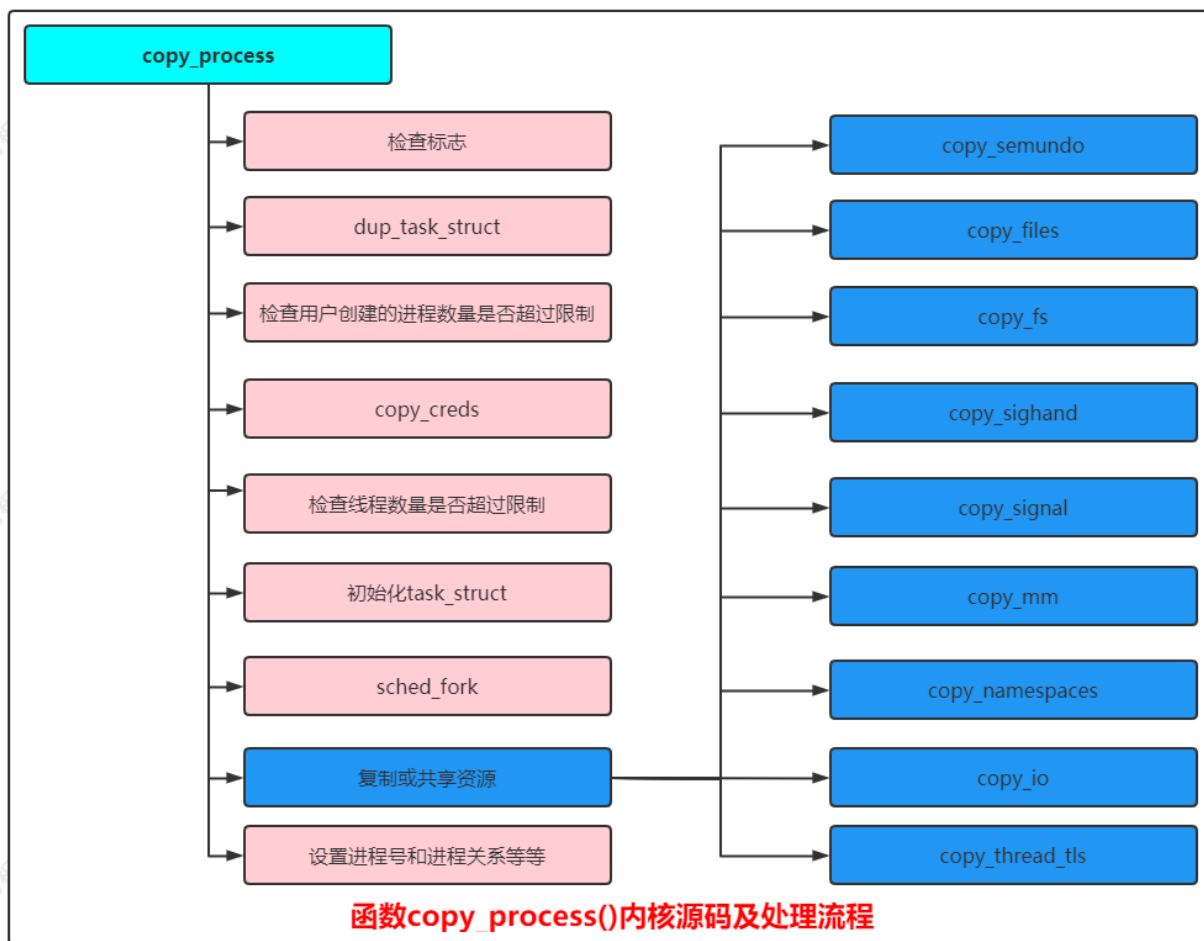
```
    free_task(p);
```

```
fork_out:
```

```
    return ERR_PTR(retval);
```

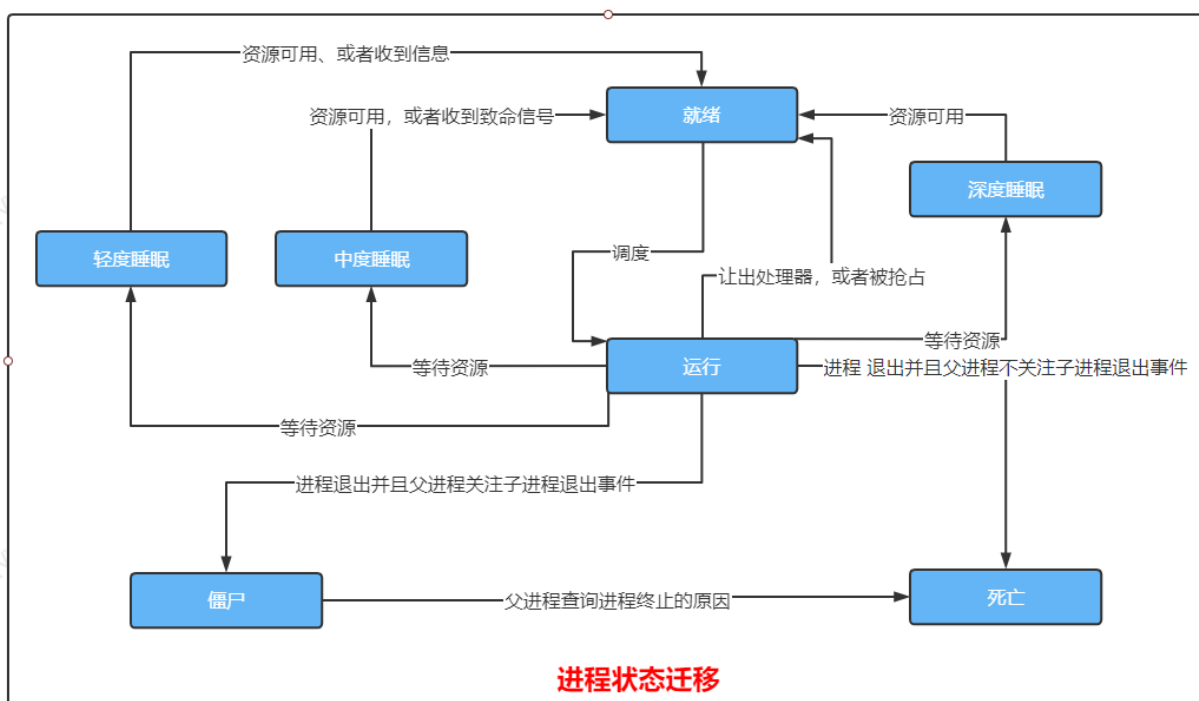
```
}
```

函数copy_process(): 创建新进程的主要工作由此函数完成, 具体处理流程如下图所示:



二、进程状态迁移

进程主要有7种状态：就绪状态、运行状态、轻度睡眠、中度睡眠、深度睡眠、僵尸状态、死亡状态，它们之间状态变迁如下：



三、调度策略及优先级

1、Linux内核支持调度策略

- 先进先出调度 (SCHED_FIFO)、轮流调度 (SCHED_RR)、限期调度策略 (SCHED_DEADLINE)采用不同的调度策略调度实时进程。
- 普通进程支持两种调度策略：标准轮流分时 (SCHED_NORMAL) 和 SCHED_BATCH调度普通的非实时进程。
- 空闲 (SCHED_IDLE) 则在系统空闲时调用idle进程。

2、进程优先级

限期进程的优先级比实时进程高，实时进程的优先级比普通进程高。

- 限制进程的优先级是-1。
- 实时进程的优先级是1-99，优先级数值越大，表示优先级越高。
- 普通进程的静态优先级是100-139，优先级值越小，表示优先级越高，可通过修改nice值改变普通进程的优先级，优先级等于120加上nice值。

在task_struct结构体中，4个成员和优先级有关如下：

```

include > linux > C sched.h > task_struct

484
485 struct task_struct {
486 #ifdef CONFIG_THREAD_INFO_IN_TASK
487     /*
488      * For reasons of header soup (see current_thread_info()), this
489      * must be the first element of task_struct.
490      */
491     struct thread_info    thread_info;
492 #endif
493     /* -1 unrunnable, 0 runnable, >0 stopped: */
494     volatile long        state;
495     void                 *stack;
496     atomic_t             usage;
497     /* Per task flags (PF_*), defined further below: */
498     unsigned int         flags;
499     unsigned int         ptrace;
500
501 #ifdef CONFIG_SMP
502     struct llist_node    wake_entry;
503     int                  on_cpu;
504 #ifdef CONFIG_THREAD_INFO_IN_TASK
505     /* Current CPU: */
506     unsigned int         cpu;
507 #endif
508     unsigned int         wakee_flips;
509     unsigned long        wakee_flip_decay_ts;
510     struct task_struct    *last_wakee;
511
512     int                  wake_cpu;
513 #endif
514     int                  on_rq;
515
516     int                  prio;
517     int                  static_prio;
518     int                  normal_prio;
519     unsigned int         rt_priority;
520
521     const struct sched_class *sched_class;

```

四、写时复制

写时复制核心思想：只有在不得不复制数据内容时才去复制数据内容。

